

COURSE
BASIC COMPUTER NETWORK

Chapter
07

TCP TRANSPORT

Editor: Nguyen Viet Ha, Ph.D.

Reference: Peter L Dordal, "An Introduction to Computer Networks," Aug. 20, 2022

March 21, 2023

Lecturer: Nguyen Viet Ha, Ph.D.

Email: nvha@hcmus.edu.vn

1. Transmission Control Protocol

❖TCP is:

➤Stream-oriented

- Application can write data in very small or very large amounts and the TCP layer will take care of appropriate packetization.

➤Connection-oriented

- Established before the beginning of any data transfer.

➤Reliable

- Correct order of delivery
- Timeout/retransmission mechanism

➤Congestion control

- TCP automatically uses the **sliding windows algorithm** to achieve throughput relatively close to the maximum available.

1

Transmission Control Protocol

2

1. Transmission Control Protocol

❖The End-to-End Principle

- It states in effect that **transport issues** are the responsibility of the **endpoints** (not the core network).

○Data corruption

- For the first, even though essentially all links on the Internet have link-layer checksums to protect against data corruption, TCP still adds its own **checksum**.

○Congestion

- TCP is today essentially the **only layer** that addresses **congestion management**.

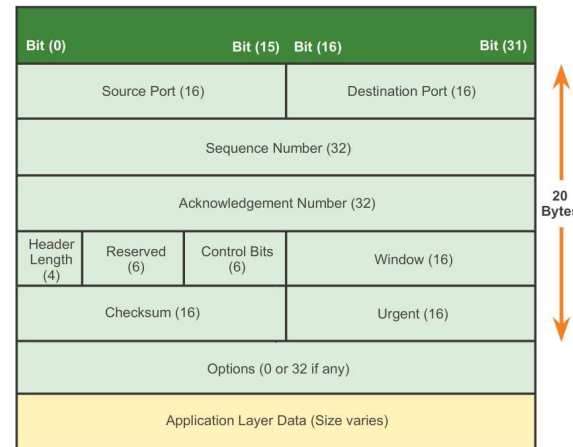
2

TCP Header

5

2. TCP Header

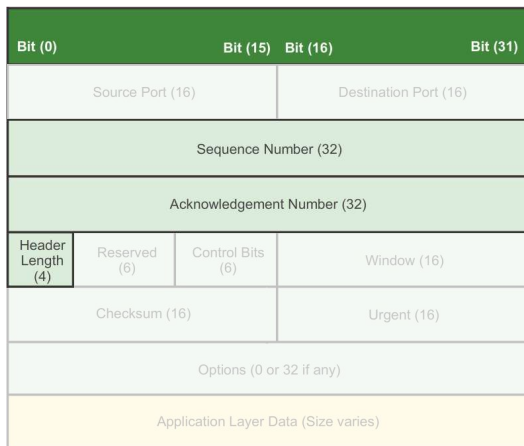
TCP Segment



❖ It is traditional to refer to the data portion of TCP packets (PDU – Packet Data Unit) as **segments**.

2. TCP Header

TCP Segment



Sequence number (32 bits) - numbering the data, at the byte level.

- The **first byte** of the **current** data payload.

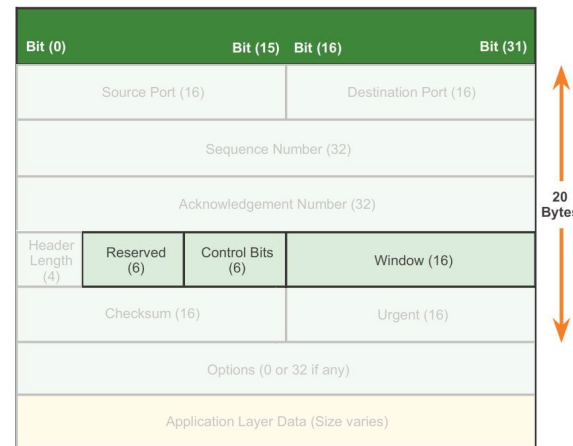
Acknowledgement number (32 bits) - Indicates the data that has been received.

- The **first byte** of the **next** data payload.

Header length (4 bits) - Indicates the length of the TCP segment header.

2. TCP Header

TCP Segment



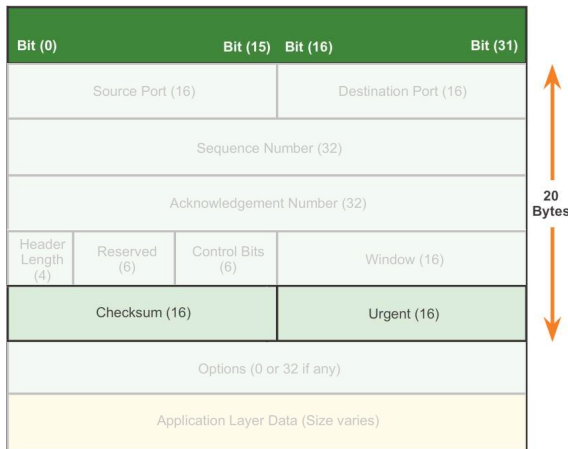
Reserved (6 bits) - is reserved for the future.

Control bits (6 bits) - Includes bit codes, or **flags**, that indicate the purpose and function of the TCP segment.

Window size (16 bits) - Indicates the number of segments that can be accepted at one time.

2. TCP Header

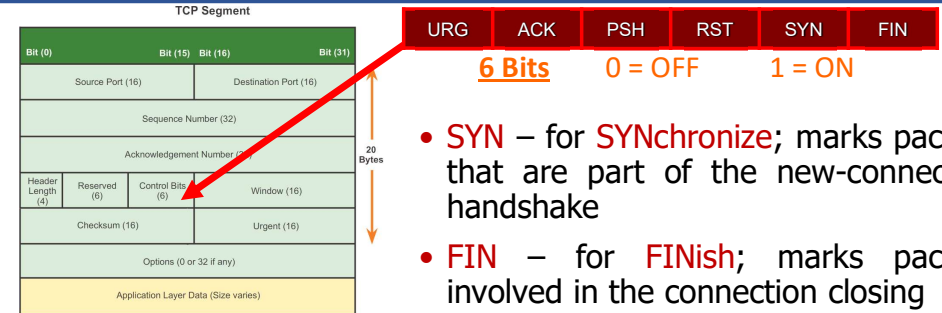
TCP Segment



Checksum (16 bits) - Used for error checking of the segment header and data.

Urgent pointer (16 bits) - Indicates if data is urgent.

2. TCP Header



- **SYN** – for **SYN**chronize; marks packets that are part of the new-connection handshake
- **FIN** – for **FIN**ish; marks packets involved in the connection closing

- **RST** – for **ReSeT**; indicates various error conditions
- **ACK** – indicates that the header **Acknowledgment** field is valid; that is, all but the first packet.
- **PSH** – for **PuSH**; marks “non-full” packets that should be delivered promptly at the far end.
- **URG** – for **URG**ent; part of a now-seldom-used mechanism for high-priority data.

2. TCP Header

❖PSH:

➤If A sends a series of **small packets** to B, then B has the option of **assembling them into a full-sized I/O buffer** before **releasing them** to the receiving application.

- However, if A sets the **PSH** bit on each packet, then B should **release each packet immediately** to the receiving application.

2. TCP Header

❖URG:

➤In **telnet** connection, A sent a large amount of data to B. Suddenly, A **wishes to abort that processing** by sending the interrupt character CNTL-C.

- Under normal conditions**, the application at B would **have to finish processing all the pending data before getting to the CNTL-C**.
- However, if the **URG bit** is set, and the **TCP header's Urgent Pointer** field **points to the CNTL-C** in the current packet, the receiving application then **skips ahead in its processing of the arriving data stream until it reaches the urgent data**.


3

TCP Connection Establishment

13

3. TCP Connection Establishment

❖ TCP connections are established via an exchange known as the **three-way handshake**.

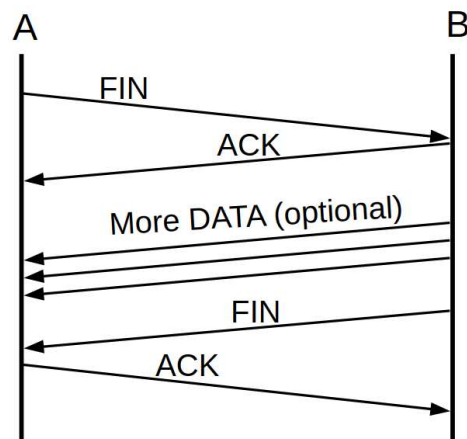


Nguyen Viet Ha, Ph.D. - Department of Telecommunications and Networks, FETEL, HCMUS, HCM-VNU

14/29

3. TCP Connection Establishment


❖ Close the connection: **two-way FIN/ACK handshakes**.



15/29

3. TCP Connection Establishment

❖ Example of a full exchange of packets in a representative connection.



	A sends	B sends
1	SYN, seq=0	
2		SYN+ACK, seq=0, ack=1 (expecting)
3	ACK, seq=1, ack=1 (ACK of SYN)	
4	"abc", seq=1, ack=1	
5		ACK, seq=1, ack=4
6	"defg", seq=4, ack=1	
7		seq=1, ack=8
8	"foobar", seq=8, ack=1	
9		seq=1, ack=14, "hello"
10	seq=14, ack=6, "goodbye"	
11,12	seq=21, ack=6, FIN	seq=6, ack=21 ;; ACK of "goodbye", crossing packets
13		seq=6, ack=22 ;; ACK of FIN
14		seq=6, ack=22, FIN
15	seq=22, ack=7 ;; ACK of FIN	

Nguyen Viet Ha, Ph.D. - Department of Telecommunications and Networks, FETEL, HCMUS, HCM-VNU

16/29

3. TCP Connection Establishment

- ❖ Each side chooses its **Initial Sequence Number (ISN)**, and sends that in its **initial SYN**.
 - All further **sequence numbers** sent are the **ISN** chosen by that side **plus** the **relative sequence number**.
- ❖ It **helps** with the allocation of a sequence number that does **not conflict** with other data bytes transmitted over a TCP connection.

3. TCP Connection Establishment

- ❖ If B had not been **LISTENing** at the port to which A sent its SYN, its response would have been **RST** ("reset"), meaning in this context "**connection refused**".
- ❖ Similarly, if A sent data to B before the SYN packet, the response would have been RST.
- ❖ **RST** can be **sent** by either side at **any time** **to abort** the connection.

4

Path MTU Discovery

4. Path MTU Discovery

- ❖ TCP connections **are more efficient** if they can keep **large packets** flowing between the endpoints.
- ❖ **Once upon a time**, TCP endpoints included just **512 bytes of data in each packet** that was **not destined for local delivery**, to **avoid fragmentation**.
- ❖ TCP endpoints **now** typically engage in **Path MTU Discovery** which almost always allows them to send larger packets.
 - Backbone ISPs are now usually able to carry 1500-byte packets.

4. Path MTU Discovery

❖The **IPv4 strategy** is to send an initial data packet with the IPv4 DONT_FRAG bit set.

➤If the **ICMP message Frag_Required/DONT_FRAG_Set** comes back, or if the packet times out, the sender **tries a smaller size**.

➤If the sender **receives a TCP ACK for the packet**, on the other hand, indicating that it made it through to the other end, it **might try a larger size**.

4. Path MTU Discovery

❖IPv6 has no DONT_FRAG bit.

❖Path **MTU Discovery over IPv6** involves the **periodic** sending of larger packets; if the **ICMPv6 message Packet Too Big** is received, a **smaller packet size** must be used.

5

TCP Flow Control

5. TCP Flow Control

❖**TCP Sliding Windows** (are measured in terms of bytes)

➤To improve throughput.

➤In the **initial three-way handshake**, each side specifies the maximum window size it is willing to accept, in the **Window Size** field of the TCP header.

○This 16-bit field can only go to **65,535 Bytes**.

▪ **Window Scale** option that can also be negotiated in the opening handshake to **increase the Window Size**.

○The window size included in the TCP header is known as the **Advertised Window Size**.

➤**TCP may either transmit a bulk stream of data**, using sliding windows fully, or it may send slowly generated interactive data.

Source Port (16)		Destination Port (16)	
Sequence Number (32)			
Acknowledgement Number (32)			
Header Length (4)	Reserved (6)	Control Bits (6)	Window (16)
Checksum (16)		Urgent (16)	

5. TCP Flow Control

❖TCP Flow Control

- It is possible for a TCP sender to send data **faster than the receiver** can process it.
 - When this happens, a TCP receiver may **reduce the advertised Window Size** value of an open connection
 - To inform the sender to switch to a **smaller window size**.

5. TCP Flow Control

❖Delayed ACKs

- Simply mean that the **ACK traffic volume** is reduced.
- Because ACKs are cumulative, **one ACK from the receiver can in principle acknowledge multiple data packets** from the sender.
- Default** number of delayed ACKs is **2**.
- The maximum **ACK delay timeout** is **500 ms**.
 - Default is **200 ms**.

6

TCP Timeout and Retransmission

6. TCP Timeout and Retransmission

- ❖When TCP sends a **packet containing user data** (this excludes ACK-only packets), it sets a timeout.
 - If that **timeout expires before the packet data is acknowledged**, it is **retransmitted**.
 - If the retransmission **loss** the sender **doubles Timeout**.
 - Retrying **5 times** as the default.

THANK YOU FOR YOUR ATTENTION

UNIVERSITY
Of SCIENCE
Est. 1941

VNUHCM



Nguyen Viet Ha, Ph.D.

Department of Telecommunications and Networks

Faculty of Electronics and Communications

University of Science, Vietnam National University, Ho Chi Minh City

Email: nvha@hcmus.edu.vn